

## Multi-label image classifier for images of musicians

Based on: <https://www.tensorflow.org/tutorials/images/classification>

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

```
Mounted at /content/drive
```

### Imports

```
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from keras import utils
from keras import preprocessing

import matplotlib.pyplot as plt
import numpy as np
import os
import PIL
import pathlib

# Load the images
data_dir = pathlib.Path('drive/MyDrive/Artists')
data_dir

PosixPath('drive/MyDrive/Artists')

michael_jackson = list(data_dir.glob('michael-jackson/*'))
print(len(michael_jackson))

50

PIL.Image.open(str(michael_jackson[0]))
```



## Data Preprocessing

```
BATCH_SIZE = 10
IMG_HEIGHT = 180
IMG_WIDTH = 180

EPOCHS = 100

train_dataset = preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(IMG_HEIGHT, IMG_WIDTH),
    batch_size=BATCH_SIZE)

class_names = train_dataset.class_names
NUM_CLASSES = len(class_names)

    Found 250 files belonging to 4 classes.
    Using 200 files for training.

val_dataset = preprocessing.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(IMG_HEIGHT, IMG_WIDTH),
    batch_size=BATCH_SIZE)

    Found 250 files belonging to 4 classes.
    Using 50 files for validation.

for image_batch, labels_batch in train_dataset:
    print(image_batch.shape)
    print(labels_batch.shape)
    break

    (10, 180, 180, 3)
    (10,)
```

```

AUTOTUNE = tf.data.AUTOTUNE
train_dataset = train_dataset.cache().shuffle(10).prefetch(buffer_size=AUTOTUNE)
val_dataset = val_dataset.cache().shuffle(10).prefetch(buffer_size=AUTOTUNE)

# Normalize
# normalization_layer = tf.keras.layers.experimental.preprocessing.Rescaling(1./255)

# Build the model
print(f'Building a {NUM_CLASSES}-label classifier.')

model = Sequential([
    layers.experimental.preprocessing.Rescaling(1./255, input_shape=(IMG_HEIGHT, IMG_WIDTH, 3)),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(128, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(NUM_CLASSES)
])

model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])

model.summary()

```

Building a 4-label classifier.  
Model: "sequential\_1"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 180, 180, 3)	0
conv2d_3 (Conv2D)	(None, 180, 180, 16)	448
max_pooling2d_3 (MaxPooling2)	(None, 90, 90, 16)	0
conv2d_4 (Conv2D)	(None, 90, 90, 32)	4640
max_pooling2d_4 (MaxPooling2)	(None, 45, 45, 32)	0
conv2d_5 (Conv2D)	(None, 45, 45, 64)	18496
max_pooling2d_5 (MaxPooling2)	(None, 22, 22, 64)	0

conv2d_6 (Conv2D)	(None, 22, 22, 128)	73856
max_pooling2d_6 (MaxPooling2D)	(None, 11, 11, 128)	0
flatten_1 (Flatten)	(None, 15488)	0
dense_2 (Dense)	(None, 128)	1982592
dense_3 (Dense)	(None, 4)	516
=====		
Total params: 2,080,548		
Trainable params: 2,080,548		
Non-trainable params: 0		

```
history = model.fit(train_dataset, validation_data=val_dataset, epochs=200)
```

```
Epoch 172/200
20/20 [=====] - 0s 10ms/step - loss: 1.3709e-07 - accuracy: 0.9999
Epoch 173/200
20/20 [=====] - 0s 10ms/step - loss: 1.3828e-07 - accuracy: 0.9999
Epoch 174/200
20/20 [=====] - 0s 10ms/step - loss: 1.3351e-07 - accuracy: 0.9999
Epoch 175/200
20/20 [=====] - 0s 9ms/step - loss: 1.3292e-07 - accuracy: 0.9999
Epoch 176/200
20/20 [=====] - 0s 10ms/step - loss: 1.2934e-07 - accuracy: 0.9999
Epoch 177/200
20/20 [=====] - 0s 9ms/step - loss: 1.2875e-07 - accuracy: 0.9999
Epoch 178/200
20/20 [=====] - 0s 9ms/step - loss: 1.2815e-07 - accuracy: 0.9999
Epoch 179/200
20/20 [=====] - 0s 9ms/step - loss: 1.2219e-07 - accuracy: 0.9999
Epoch 180/200
20/20 [=====] - 0s 10ms/step - loss: 1.2457e-07 - accuracy: 0.9999
Epoch 181/200
20/20 [=====] - 0s 10ms/step - loss: 1.2219e-07 - accuracy: 0.9999
Epoch 182/200
20/20 [=====] - 0s 9ms/step - loss: 1.1861e-07 - accuracy: 0.9999
Epoch 183/200
20/20 [=====] - 0s 10ms/step - loss: 1.1563e-07 - accuracy: 0.9999
Epoch 184/200
20/20 [=====] - 0s 9ms/step - loss: 1.1683e-07 - accuracy: 0.9999
Epoch 185/200
20/20 [=====] - 0s 10ms/step - loss: 1.1504e-07 - accuracy: 0.9999
Epoch 186/200
20/20 [=====] - 0s 10ms/step - loss: 1.1325e-07 - accuracy: 0.9999
Epoch 187/200
20/20 [=====] - 0s 10ms/step - loss: 1.1027e-07 - accuracy: 0.9999
Epoch 188/200
20/20 [=====] - 0s 10ms/step - loss: 1.1206e-07 - accuracy: 0.9999
Epoch 189/200
20/20 [=====] - 0s 10ms/step - loss: 1.0848e-07 - accuracy: 0.9999
Epoch 190/200
20/20 [=====] - 0s 10ms/step - loss: 1.0848e-07 - accuracy: 0.9999
Epoch 191/200
20/20 [=====] - 0s 10ms/step - loss: 1.0669e-07 - accuracy: 0.9999
```

```

Epoch 192/200
20/20 [=====] - 0s 10ms/step - loss: 1.0431e-07 - accuracy:
Epoch 193/200
20/20 [=====] - 0s 10ms/step - loss: 1.0490e-07 - accuracy:
Epoch 194/200
20/20 [=====] - 0s 9ms/step - loss: 1.0312e-07 - accuracy:
Epoch 195/200
20/20 [=====] - 0s 10ms/step - loss: 1.0133e-07 - accuracy:
Epoch 196/200
20/20 [=====] - 0s 10ms/step - loss: 9.7752e-08 - accuracy:
Epoch 197/200
20/20 [=====] - 0s 10ms/step - loss: 9.8944e-08 - accuracy:
Epoch 198/200
20/20 [=====] - 0s 9ms/step - loss: 9.9540e-08 - accuracy:
Epoch 199/200
20/20 [=====] - 0s 10ms/step - loss: 9.5963e-08 - accuracy:
Epoch 200/200
20/20 [=====] - 0s 10ms/step - loss: 9.6560e-08 - accuracy:

```

```
# Get confusion matrix on the val dataset
```

```
y_pred = model.predict(val_dataset)
```

```
# Build confusion matrix
```

```
predicted_categories = np.argmax(y_pred, axis=1)
```

```
actual_categories = tf.concat([y for x, y in val_dataset], axis=0)
```

```
print(f'Predicted: {predicted_categories}')
```

```
print(f'Actual: {actual_categories}')
```

```

Predicted: [2 0 3 3 3 1 3 0 3 3 3 3 3 3 3 3 2 1 3 1 3 3 0 3 3 3 1 0 1 2 3 2 0 3
 1 3 3 1 0 2 2 1 3 0 3 0 3]
Actual: [3 2 1 1 3 0 3 3 3 1 0 1 2 3 2 1 3 2 3 1 2 0 3 3 2 2 2 1 1 1 1 0 2 2 3 3 0
 3 0 3 1 1 2 1 3 1 3 3 3]

```

```
cm = tf.math.confusion_matrix(actual_categories, predicted_categories, num_classes=4)
```

```
# This code is from:
```

```
# https://scikit-learn.org/0.18/auto\_examples/model\_selection/plot\_confusion\_matrix.html
```

```
import matplotlib.pyplot as plt
```

```
from sklearn.metrics import confusion_matrix
```

```
import itertools
```

```

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):

```

```
    """
```

```

    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

```

```
    """
```

```

plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

print(cm)

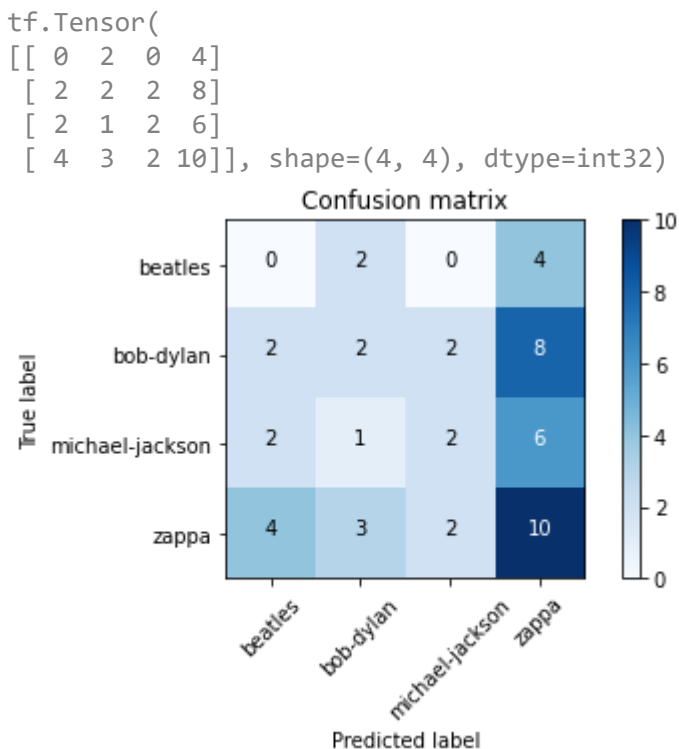
cm = cm.numpy()
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, cm[i, j],
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

plt.figure()
plot_confusion_matrix(cm, classes=class_names, normalize=True, title='Confusion matrix')

plt.show()

```



```

def predict(model, url):
    predict_url = url
    predict_path = tf.keras.utils.get_file('predict', origin=predict_url)
    print(predict_path)

```

```

img = keras.preprocessing.image.load_img(
    predict_path, target_size=(IMG_HEIGHT, IMG_WIDTH)
)

img_array = []
img_array = keras.preprocessing.image.img_to_array(img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

predictions = model.predict(img_array)
score = tf.nn.softmax(predictions[0])

print(
    "This image most likely belongs to {} with a {:.2f} percent confidence."
    .format(class_names[np.argmax(score)], 100 * np.max(score))
)

plt.imshow(img)

sample_urls = [
    "https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcSjKvc7_8SLeHlsOrLT6-Lpnpto00LDmCo_kg",
    "https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcQCQ6Fk1AbBcRczijclpOKswAm3te5vG7Mfng",
    "https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcTt3tTCLgunkc4BAP4PMsWhUxaGkSxtA6BMFv",
    "https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcTcsXPkjxp4FMKniRZKlewCz30jYrML5sfDJg"
]

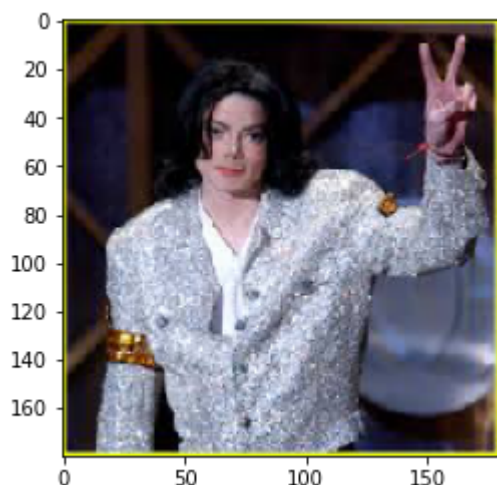
# model.save(f'saved_models/classifier')

re_model = keras.models.load_model(f'saved_models/classifier')
predict(re_model, sample_urls[1])

```

/root/.keras/datasets/predict

This image most likely belongs to zappa with a 100.00 percent confidence.



---

✓ 0s completed at 10:17 PM

